

In [1]:

```
import pandas as pd
import warnings
warnings.filterwarnings('ignore')
```

## Step 1 : Reading and Understanding the data

In [2]:

```
master_df = pd.read_csv("CarPrice_Assignment.csv")
master_df.set_index('car_ID', inplace = True)
master_df.rename(columns = {"symboling" : "Insuranceriskfactor"}, inplace = True)
```

In [3]:

```
master_df.head()
```

Out[3]:

	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drivewhl
car_ID							
1	3	alfa-romero giulia	gas	std	two	convertible	
2	3	alfa-romero stelvio	gas	std	two	convertible	
3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	
4	2	audi 100 ls	gas	std	four	sedan	
5	2	audi 100ls	gas	std	four	sedan	

5 rows × 25 columns

In [4]:

```
master_df.dtypes
```

Out[4]:

```

Insuranceriskfactor    int64
CarName                 object
fueltype               object
aspiration             object
doornumber             object
carbody               object
drivewheel            object
engineloation         object
wheelbase             float64
carlength             float64
carwidth              float64
carheight             float64
curbweight            int64
enginetype            object
cylindernumber        object
enginesize            int64
fuelsystem            object
boreratio             float64
stroke               float64
compressionratio      float64
horsepower            int64
peakrpm              int64
citympg              int64
highwaympg           int64
price                float64
dtype: object

```

In [5]:

```
master_df.describe()
```

Out[5]:

	Insuranceriskfactor	wheelbase	carlength	carwidth	carheight	curbweight	eng
<b>count</b>	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205
<b>mean</b>	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126
<b>std</b>	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41
<b>min</b>	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61
<b>25%</b>	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97
<b>50%</b>	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120
<b>75%</b>	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141
<b>max</b>	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326

## Step 2 : Visualizing the data

## Visualizing the Numerical columns

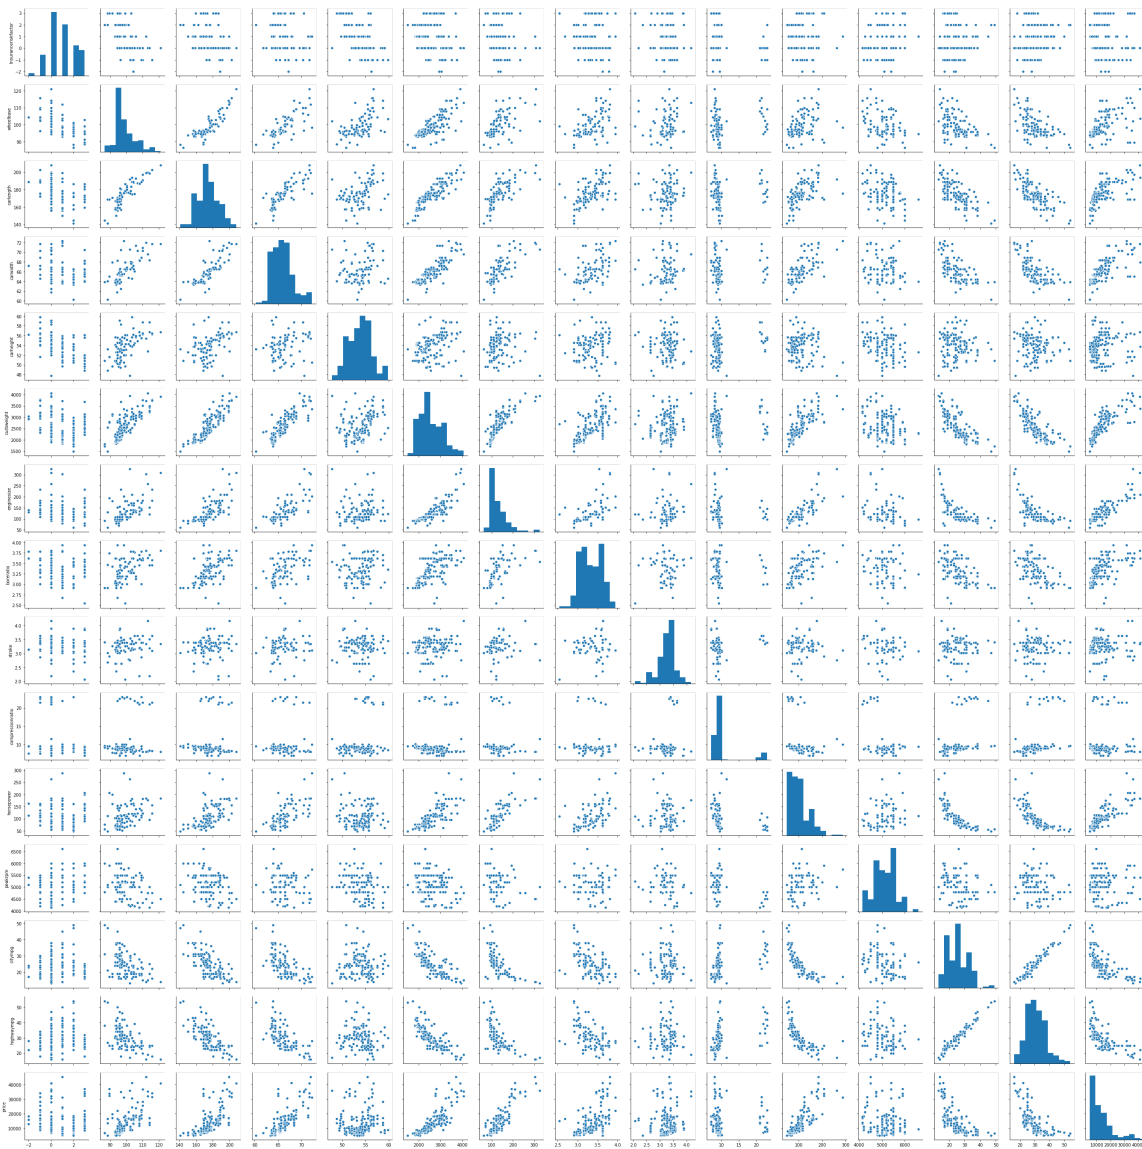
In [69]:

```
import matplotlib.pyplot as plt
import seaborn as sns

numerics = ['int16', 'int32', 'int64', 'float16', 'float32', 'float64']
num_columns_list = list(master_df.select_dtypes(include=numerics).columns)

plt.figure(figsize = (20,12))
sns.pairplot(master_df[num_columns_list])
plt.show()
```

<Figure size 1440x864 with 0 Axes>



## Visualizing the categorical variables

### Restructuring CarName to have on ly the company name

In [7]:

```
def strip_car_model(car_name):
    company_name = car_name.split()
    return company_name[0]

master_df.CarName = master_df.CarName.apply(strip_car_model)
master_df.head()
```

Out[7]:

	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
car_ID							
1	3	alfa-romero	gas	std	two	convertible	rw
2	3	alfa-romero	gas	std	two	convertible	rw
3	1	alfa-romero	gas	std	two	hatchback	rw
4	2	audi	gas	std	four	sedan	fw
5	2	audi	gas	std	four	sedan	4w

5 rows × 25 columns

In [8]:


```
non_num_columns_list = list(master_df.select_dtypes(exclude=numerics).columns)

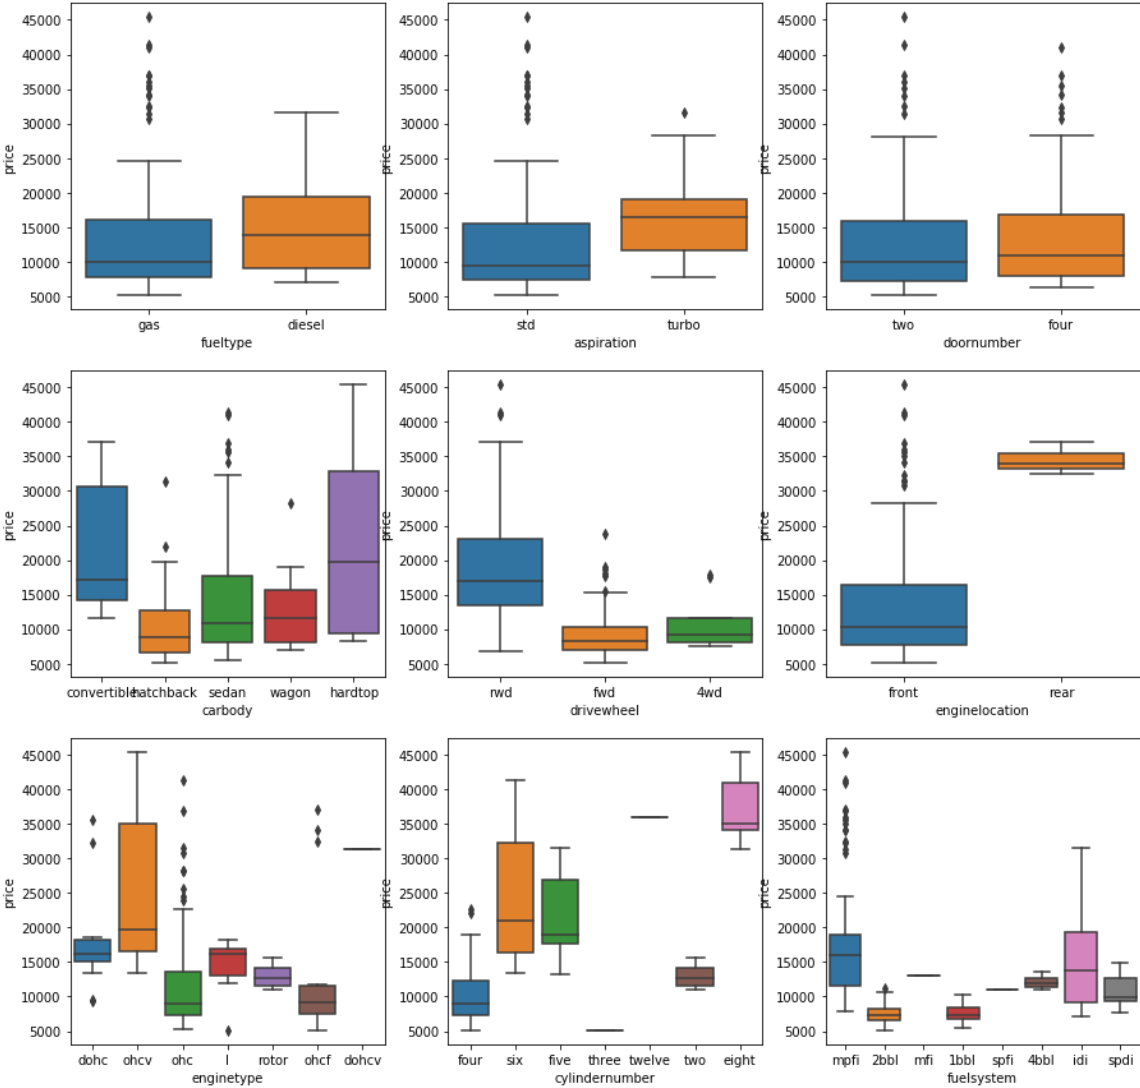
## Visualize the CarName first
sns.boxplot(x= 'CarName', y = 'price', data = master_df)
plt.xticks(rotation='vertical')
plt.show()

subplot_cnt = 1
plt.figure(figsize = (15,15))

for each_cat_var in non_num_columns_list:
    if each_cat_var != 'CarName':
        plt.subplot(3,3,subplot_cnt)
        sns.boxplot(x= each_cat_var, y = 'price', data = master_df)
        subplot_cnt = subplot_cnt+1

plt.show()
```







In [9]:

```
## Visualize the carbody and driveweheel impact on the price
plt.figure(figsize = (10, 5))
sns.boxplot(x = 'carbody', y = 'price', hue = 'drivewheel', data = master_df)
plt.show()

## Visualize the carbody and enginelocation impact on the price
plt.figure(figsize = (10, 5))
sns.boxplot(x = 'carbody', y = 'price', hue = 'enginelocation', data = master_df)
plt.show()
```





### Step 3 : Data Preparation

In [10]:

```

import sklearn
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder

## Create a copy of df
master_df_le = master_df.copy()

## Apply label encoder on all the categorical variables
le = preprocessing.LabelEncoder()

for each_item in non_num_columns_list:
    le.fit(master_df_le[each_item])
    master_df_le[each_item] = le.transform(master_df_le[each_item])

master_df_le.head()

```

Out[10]:

	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
car_ID							
1	3	1	1	0	1	0	2
2	3	1	1	0	1	0	2
3	1	1	1	0	1	2	2
4	2	2	1	0	0	3	1
5	2	2	1	0	0	3	0

5 rows × 25 columns

## Step 4: Splitting the training and test data

In [11]:

```
import numpy as np
from sklearn.model_selection import train_test_split

np.random.seed(0)
master_df_train, master_df_test = train_test_split(master_df_le, train_size = 0.7, test_size = 0.3, random_state = 100)
master_df_train.head()
```

Out[11]:

	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
car_ID							
123	1	16	1	0	0	3	1
126	3	18	1	0	1	2	2
167	1	22	1	0	1	2	2
2	3	1	1	0	1	0	2
200	-1	26	1	1	0	4	2

5 rows × 25 columns

## Re-sclaing the parameters

In [12]:

```

## Scale and transform all the numerical values
## especially the car dimentions, wheelbase and price are on very hig scale in comparis
on to the
## let us use minmaxmethod - as it helps normalize the outliers

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
master_df_train[num_columns_list] = scaler.fit_transform(master_df_train[num_columns_li
st])
master_df_train[non_num_columns_list] = scaler.fit_transform(master_df_train[non_num_co
lums_list])
master_df_train.head()

```

Out[12]:

	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
car_ID							
123	0.6	0.592593	1.0	0.0	0.0	0.75	0.5
126	1.0	0.666667	1.0	0.0	1.0	0.50	1.0
167	0.6	0.814815	1.0	0.0	1.0	0.50	1.0
2	1.0	0.037037	1.0	0.0	1.0	0.00	1.0
200	0.2	0.962963	1.0	1.0	0.0	1.00	1.0

5 rows × 25 columns

In [13]:

```
master_df_train.describe()
```

Out[13]:


	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drive
<b>count</b>	143.000000	143.000000	143.000000	143.000000	143.000000	143.000000	143
<b>mean</b>	0.559441	0.517742	0.909091	0.181818	0.440559	0.666084	0
<b>std</b>	0.239200	0.276478	0.288490	0.387050	0.498199	0.209678	0
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0
<b>25%</b>	0.400000	0.259259	1.000000	0.000000	0.000000	0.500000	0
<b>50%</b>	0.600000	0.518519	1.000000	0.000000	0.000000	0.750000	0
<b>75%</b>	0.600000	0.777778	1.000000	0.000000	1.000000	0.750000	1
<b>max</b>	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1

8 rows × 25 columns

**Calculating the correlation between different parameters of the dataframe.**

In [14]:

```
plt.figure(figsize = (20, 20))  
sns.heatmap(master_df_train.corr(), annot = True, cmap="YlGnBu")  
plt.show()
```



## Step 5: Building a linear model

In [15]:

```
## Prepare the X and y train data

cols_list = list(master_df_train.columns)
cols_list.remove('price')

y_train = master_df_train['price']
X_train = master_df_train[cols_list]
```

In [16]:

```
y_train.head()
```

Out[16]:

```
car_ID
123    0.068818
126    0.466890
167    0.122110
2      0.314446
200    0.382131
Name: price, dtype: float64
```

In [17]:

```
X_train.head()
```

Out[17]:

	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
car_ID							
123	0.6	0.592593	1.0	0.0	0.0	0.75	0.5
126	1.0	0.666667	1.0	0.0	1.0	0.50	1.0
167	0.6	0.814815	1.0	0.0	1.0	0.50	1.0
2	1.0	0.037037	1.0	0.0	1.0	0.00	1.0
200	0.2	0.962963	1.0	1.0	0.0	1.00	1.0

5 rows × 24 columns

## Model 1 : Adding all the parameters (predictors) to build the model



In [18]:

```
import statsmodels.api as sm
X_train_lm = sm.add_constant(X_train)
lr_1 = sm.OLS(y_train, X_train_lm).fit()
lr_1.summary()
```

Out[18]:

## OLS Regression Results

Dep. Variable:	price		R-squared:		0.886		
Model:	OLS		Adj. R-squared:		0.863		
Method:	Least Squares		F-statistic:		38.15		
Date:	Wed, 02 Oct 2019		Prob (F-statistic):		1.45e-44		
Time:	10:34:06		Log-Likelihood:		172.11		
No. Observations:	143		AIC:		-294.2		
Df Residuals:	118		BIC:		-220.2		
Df Model:	24						
Covariance Type:	nonrobust						
	coef	std err	t	P> t	[0.025	0.975]	
const	-0.2899	0.302	-0.959	0.340	-0.889	0.309	
Insuranceriskfactor	0.0232	0.049	0.475	0.635	-0.073	0.120	
CarName	-0.1349	0.031	-4.397	0.000	-0.196	-0.074	
fueltype	0.1255	0.253	0.496	0.621	-0.376	0.627	
aspiration	0.0115	0.031	0.370	0.712	-0.050	0.073	
doornumber	-0.0134	0.023	-0.597	0.552	-0.058	0.031	
carbody	-0.0628	0.055	-1.150	0.252	-0.171	0.045	
drivewheel	0.0525	0.038	1.376	0.172	-0.023	0.128	
enginelocation	0.3261	0.095	3.449	0.001	0.139	0.513	
wheelbase	0.0696	0.111	0.627	0.532	-0.150	0.290	
carlength	-0.0765	0.121	-0.631	0.529	-0.317	0.164	
carwidth	0.2379	0.126	1.886	0.062	-0.012	0.488	
carheight	0.0897	0.055	1.638	0.104	-0.019	0.198	
curbweight	0.2757	0.146	1.893	0.061	-0.013	0.564	
enginetype	0.0280	0.049	0.577	0.565	-0.068	0.124	
cylindernumber	0.0449	0.072	0.624	0.534	-0.098	0.187	
engineize	0.4803	0.160	3.008	0.003	0.164	0.796	
fuelsystem	0.0302	0.035	0.872	0.385	-0.038	0.099	
boreratio	-0.0124	0.054	-0.229	0.819	-0.119	0.095	
stroke	-0.1081	0.061	-1.774	0.079	-0.229	0.013	
compressionratio	0.1755	0.292	0.602	0.548	-0.402	0.753	
horsepower	0.1702	0.157	1.087	0.279	-0.140	0.480	
peakrpm	0.0537	0.056	0.968	0.335	-0.056	0.164	
citympg	-0.0530	0.211	-0.251	0.802	-0.471	0.365	
highwaympg	0.0851	0.199	0.428	0.669	-0.308	0.479	
Omnibus:	44.657	Durbin-Watson:	1.814				

<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	245.736
<b>Skew:</b>	0.933	<b>Prob(JB):</b>	4.36e-54
<b>Kurtosis:</b>	9.145	<b>Cond. No.</b>	184.

**Warnings:**

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [19]:

```

## Calculate the VIF
from statsmodels.stats.outliers_influence import variance_inflation_factor

vif = pd.DataFrame()
vif['Features'] = X_train.columns
vif['VIF'] = [variance_inflation_factor(X_train.values, i) for i in range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

Out[19]:

	Features	VIF
2	fueltype	157.95
23	highwaympg	155.34
22	citympg	152.24
9	carlength	104.38
12	curbweight	94.01
10	carwidth	83.98
8	wheelbase	56.21
15	enginesize	45.88
20	horsepower	43.08
19	compressionratio	32.77
5	carbody	32.47
11	carheight	20.26
18	stroke	20.26
0	Insuranceriskfactor	18.48
17	boreratio	17.16
14	cylindernumber	16.77
6	drivewheel	16.13
13	enginetype	14.34
21	peakrpm	13.53
16	fuelsystem	9.69
1	CarName	6.78
4	doornumber	4.99
3	aspiration	2.81
7	enginelocation	1.40

## **Model 2 - Dropping the fueltype predictor as it has the highest P value (insignificane)**

In [20]:

```
X = X_train.drop('fueltype',1)
X_train_lm = sm.add_constant(X)
lr_2 = sm.OLS(y_train, X_train_lm).fit()
lr_2.summary()
```

Out[20]:

## OLS Regression Results

Dep. Variable:	price		R-squared:		0.886		
Model:	OLS		Adj. R-squared:		0.863		
Method:	Least Squares		F-statistic:		40.05		
Date:	Wed, 02 Oct 2019		Prob (F-statistic):		2.54e-45		
Time:	10:34:06		Log-Likelihood:		171.96		
No. Observations:	143		AIC:		-295.9		
Df Residuals:	119		BIC:		-224.8		
Df Model:	23						
Covariance Type:	nonrobust						
	coef	std err	t	P> t	[0.025	0.975]	
const	-0.1494	0.105	-1.423	0.157	-0.357	0.058	
Insuranceriskfactor	0.0202	0.048	0.419	0.676	-0.075	0.116	
CarName	-0.1302	0.029	-4.481	0.000	-0.188	-0.073	
aspiration	0.0035	0.026	0.132	0.895	-0.049	0.056	
doornumber	-0.0131	0.022	-0.585	0.560	-0.058	0.031	
carbody	-0.0622	0.054	-1.144	0.255	-0.170	0.045	
drivewheel	0.0548	0.038	1.452	0.149	-0.020	0.129	
enginelocation	0.3239	0.094	3.440	0.001	0.137	0.510	
wheelbase	0.0581	0.108	0.536	0.593	-0.156	0.272	
carlength	-0.0686	0.120	-0.573	0.568	-0.306	0.169	
carwidth	0.2515	0.123	2.049	0.043	0.008	0.495	
carheight	0.0887	0.055	1.626	0.107	-0.019	0.197	
curbweight	0.2624	0.143	1.839	0.068	-0.020	0.545	
enginetype	0.0253	0.048	0.526	0.600	-0.070	0.121	
cylindernumber	0.0550	0.069	0.800	0.425	-0.081	0.191	
engineize	0.4927	0.157	3.135	0.002	0.181	0.804	
fuelsystem	0.0264	0.034	0.785	0.434	-0.040	0.093	
boreratio	-0.0194	0.052	-0.372	0.710	-0.122	0.084	
stroke	-0.1207	0.055	-2.186	0.031	-0.230	-0.011	
compressionratio	0.0330	0.049	0.679	0.498	-0.063	0.129	
horsepower	0.1762	0.156	1.132	0.260	-0.132	0.484	
peakrpm	0.0599	0.054	1.112	0.268	-0.047	0.167	
citympg	-0.0492	0.210	-0.234	0.816	-0.466	0.367	
highwaympg	0.0885	0.198	0.447	0.656	-0.303	0.480	
Omnibus:	40.492	Durbin-Watson:		1.821			
Prob(Omnibus):	0.000	Jarque-Bera (JB):		199.345			

**Skew:** 0.859      **Prob(JB):** 5.16e-44  
**Kurtosis:** 8.523      **Cond. No.** 105.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [21]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[21]:

	Features	VIF
22	highwaympg	153.12
21	citympg	143.50
8	carlength	98.53
11	curbweight	92.30
9	carwidth	83.72
7	wheelbase	55.66
14	enginesize	45.58
19	horsepower	43.02
4	carbody	31.89
17	stroke	20.17
10	carheight	19.69
16	boreratio	17.01
0	Insuranceriskfactor	16.72
13	cylindernumber	16.19
5	drivewheel	15.74
12	enginetype	12.85
20	peakrpm	11.61
15	fuelsystem	9.69
1	CarName	6.44
3	doornumber	4.97
18	compressionratio	4.17
2	aspiration	2.79
6	enginelocation	1.39



In [ ]:

**Model 3 - Dropping the fueltype predictor as it has the highest P value (insignificane)**

In [22]:

```
X = X.drop('Insuranceriskfactor',1)
X_train_lm = sm.add_constant(X)
lr_3 = sm.OLS(y_train, X_train_lm).fit()
lr_3.summary()
```

Out[22]:

## OLS Regression Results

Dep. Variable:	price		R-squared:		0.885	
Model:	OLS		Adj. R-squared:		0.864	
Method:	Least Squares		F-statistic:		42.15	
Date:	Wed, 02 Oct 2019		Prob (F-statistic):		4.20e-46	
Time:	10:34:07		Log-Likelihood:		171.86	
No. Observations:	143		AIC:		-297.7	
Df Residuals:	120		BIC:		-229.6	
Df Model:	22					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-0.1332	0.097	-1.369	0.174	-0.326	0.059
CarName	-0.1312	0.029	-4.548	0.000	-0.188	-0.074
aspiration	0.0034	0.026	0.130	0.897	-0.049	0.055
doornumber	-0.0094	0.021	-0.458	0.648	-0.050	0.031
carbody	-0.0656	0.054	-1.222	0.224	-0.172	0.041
drivewheel	0.0544	0.038	1.447	0.150	-0.020	0.129
enginelocation	0.3287	0.093	3.531	0.001	0.144	0.513
wheelbase	0.0405	0.100	0.407	0.685	-0.156	0.238
carlength	-0.0670	0.119	-0.561	0.576	-0.303	0.169
carwidth	0.2628	0.119	2.201	0.030	0.026	0.499
carheight	0.0877	0.054	1.615	0.109	-0.020	0.195
curbweight	0.2681	0.142	1.894	0.061	-0.012	0.548
enginetype	0.0210	0.047	0.448	0.655	-0.072	0.114
cylindernumber	0.0610	0.067	0.909	0.365	-0.072	0.194
enginesize	0.4882	0.156	3.124	0.002	0.179	0.798
fuelsystem	0.0287	0.033	0.868	0.387	-0.037	0.094
boreratio	-0.0185	0.052	-0.357	0.722	-0.121	0.084
stroke	-0.1209	0.055	-2.197	0.030	-0.230	-0.012
compressionratio	0.0323	0.048	0.669	0.505	-0.063	0.128
horsepower	0.1623	0.152	1.071	0.286	-0.138	0.462
peakrpm	0.0605	0.054	1.127	0.262	-0.046	0.167
citympg	-0.0592	0.208	-0.284	0.777	-0.472	0.353
highwaympg	0.0934	0.197	0.474	0.636	-0.296	0.483
Omnibus:	38.950	Durbin-Watson:		1.816		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		181.348		
Skew:	0.838	Prob(JB):		4.18e-40		

**Kurtosis:** 8.256**Cond. No.**

101.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [23]:

```

vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

Out[23]:

	Features	VIF
21	highwaympg	151.28
20	citympg	143.39
7	carlength	97.12
10	curbweight	88.96
8	carwidth	79.46
6	wheelbase	45.55
13	enginesize	45.40
18	horsepower	40.86
3	carbody	31.54
16	stroke	20.01
9	carheight	19.67
15	boreratio	16.75
4	drivewheel	15.73
12	cylindernumber	15.11
11	enginetype	12.76
19	peakrpm	11.29
14	fuelsystem	9.29
0	CarName	6.43
17	compressionratio	4.00
2	doornumber	3.97
1	aspiration	2.78
5	enginelocation	1.37

In [ ]:

## **Model 4 - Dropping the carlength predictor as it has the highest P value and VIF (insignificane)**

In [24]:

```
X = X.drop('carlength',1)
X_train_lm = sm.add_constant(X)
lr_4 = sm.OLS(y_train, X_train_lm).fit()
lr_4.summary()
```

Out[24]:

## OLS Regression Results

Dep. Variable:	price		R-squared:		0.885	
Model:	OLS		Adj. R-squared:		0.865	
Method:	Least Squares		F-statistic:		44.40	
Date:	Wed, 02 Oct 2019		Prob (F-statistic):		7.24e-47	
Time:	10:34:07		Log-Likelihood:		171.67	
No. Observations:	143		AIC:		-299.3	
Df Residuals:	121		BIC:		-234.2	
Df Model:	21					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-0.1460	0.094	-1.548	0.124	-0.333	0.041
CarName	-0.1313	0.029	-4.566	0.000	-0.188	-0.074
aspiration	0.0057	0.026	0.222	0.825	-0.045	0.057
doornumber	-0.0081	0.020	-0.397	0.692	-0.048	0.032
carbody	-0.0752	0.051	-1.484	0.140	-0.176	0.025
drivewheel	0.0536	0.037	1.430	0.155	-0.021	0.128
enginelocation	0.3252	0.093	3.511	0.001	0.142	0.509
wheelbase	0.0294	0.097	0.302	0.763	-0.163	0.222
carwidth	0.2381	0.111	2.152	0.033	0.019	0.457
carheight	0.0835	0.054	1.557	0.122	-0.023	0.190
curbweight	0.2565	0.140	1.837	0.069	-0.020	0.533
enginetype	0.0282	0.045	0.628	0.531	-0.061	0.117
cylindernumber	0.0556	0.066	0.840	0.403	-0.075	0.187
engineize	0.4787	0.155	3.090	0.002	0.172	0.785
fuelsystem	0.0239	0.032	0.751	0.454	-0.039	0.087
boreratio	-0.0228	0.051	-0.446	0.656	-0.124	0.078
stroke	-0.1203	0.055	-2.192	0.030	-0.229	-0.012
compressionratio	0.0306	0.048	0.637	0.526	-0.065	0.126
horsepower	0.1826	0.147	1.244	0.216	-0.108	0.473
peakrpm	0.0622	0.053	1.163	0.247	-0.044	0.168
citympg	-0.0352	0.203	-0.173	0.863	-0.438	0.367
highwaympg	0.0886	0.196	0.451	0.652	-0.300	0.477
Omnibus:	35.678	Durbin-Watson:	1.811			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	153.022			
Skew:	0.776	Prob(JB):	5.91e-34			
Kurtosis:	7.824	Cond. No.	96.9			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [25]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[25]:

	Features	VIF
20	highwaympg	150.28
19	citympg	139.60
9	curbweight	84.59
7	carwidth	67.10
12	enginesize	44.83
6	wheelbase	44.36
17	horsepower	38.03
3	carbody	27.64
15	stroke	19.98
8	carheight	18.95
14	boreratio	15.99
4	drivewheel	15.65
11	cylindernumber	14.54
10	enginetype	12.17
18	peakrpm	11.28
13	fuelsystem	8.46
0	CarName	6.42
16	compressionratio	3.99
2	doornumber	3.95
1	aspiration	2.73
5	enginelocation	1.37

In [ ]:

**Model 5 - Dropping the fuelsystem predictor as it has the highest P value (insignificane)**



In [26]:

```
X = X.drop('fuelsystem',1)
X_train_lm = sm.add_constant(X)
lr_5 = sm.OLS(y_train, X_train_lm).fit()
lr_5.summary()
```

Out[26]:

## OLS Regression Results

Dep. Variable:	price		R-squared:		0.885	
Model:	OLS		Adj. R-squared:		0.866	
Method:	Least Squares		F-statistic:		46.76	
Date:	Wed, 02 Oct 2019		Prob (F-statistic):		1.37e-47	
Time:	10:34:07		Log-Likelihood:		171.34	
No. Observations:	143		AIC:		-300.7	
Df Residuals:	122		BIC:		-238.5	
Df Model:	20					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-0.1310	0.092	-1.424	0.157	-0.313	0.051
CarName	-0.1258	0.028	-4.533	0.000	-0.181	-0.071
aspiration	0.0093	0.025	0.365	0.716	-0.041	0.059
doornumber	-0.0070	0.020	-0.345	0.730	-0.047	0.033
carbody	-0.0794	0.050	-1.580	0.117	-0.179	0.020
drivewheel	0.0537	0.037	1.435	0.154	-0.020	0.128
enginelocation	0.3221	0.092	3.487	0.001	0.139	0.505
wheelbase	0.0407	0.096	0.425	0.672	-0.149	0.231
carwidth	0.2275	0.110	2.076	0.040	0.011	0.444
carheight	0.0814	0.053	1.522	0.131	-0.024	0.187
curbweight	0.2621	0.139	1.882	0.062	-0.014	0.538
enginetype	0.0275	0.045	0.614	0.540	-0.061	0.116
cylindernumber	0.0499	0.066	0.761	0.448	-0.080	0.180
engineize	0.4754	0.155	3.075	0.003	0.169	0.781
boreratio	-0.0241	0.051	-0.472	0.637	-0.125	0.077
stroke	-0.1192	0.055	-2.177	0.031	-0.228	-0.011
compressionratio	0.0293	0.048	0.610	0.543	-0.066	0.124
horsepower	0.2009	0.144	1.390	0.167	-0.085	0.487
peakrpm	0.0598	0.053	1.122	0.264	-0.046	0.165
citympg	-0.0725	0.197	-0.369	0.713	-0.462	0.317
highwaympg	0.1113	0.193	0.575	0.566	-0.272	0.494
Omnibus:	35.529	Durbin-Watson:	1.814			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	147.274			
Skew:	0.786	Prob(JB):	1.05e-32			
Kurtosis:	7.717	Cond. No.	92.9			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [27]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[27]:

	Features	VIF
19	highwaympg	144.56
18	citympg	133.13
9	curbweight	83.21
7	carwidth	66.48
12	enginesize	44.80
6	wheelbase	43.65
16	horsepower	37.27
3	carbody	27.48
14	stroke	19.84
8	carheight	18.95
13	boreratio	15.98
4	drivewheel	15.64
11	cylindernumber	14.47
10	enginetype	12.15
17	peakrpm	11.28
0	CarName	5.89
15	compressionratio	3.91
2	doornumber	3.89
1	aspiration	2.61
5	enginelocation	1.37

**Model 6 - Dropping the cylindernumber predictor as it has the highest P value (insignificane)**

In [28]:

```
X = X.drop('cylindernumber',1)
X_train_lm = sm.add_constant(X)
lr_6 = sm.OLS(y_train, X_train_lm).fit()
lr_6.summary()
```

Out[28]:

## OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.884
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.866
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	49.36
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	2.55e-48
<b>Time:</b>	10:34:07	<b>Log-Likelihood:</b>	171.00
<b>No. Observations:</b>	143	<b>AIC:</b>	-302.0
<b>Df Residuals:</b>	123	<b>BIC:</b>	-242.7
<b>Df Model:</b>	19		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.1165	0.090	-1.297	0.197	-0.294	0.061
<b>CarName</b>	-0.1239	0.028	-4.491	0.000	-0.179	-0.069
<b>aspiration</b>	0.0028	0.024	0.118	0.906	-0.044	0.050
<b>doornumber</b>	-0.0067	0.020	-0.330	0.742	-0.047	0.033
<b>carbody</b>	-0.0746	0.050	-1.499	0.136	-0.173	0.024
<b>drivewheel</b>	0.0609	0.036	1.688	0.094	-0.011	0.132
<b>enginelocation</b>	0.3205	0.092	3.477	0.001	0.138	0.503
<b>wheelbase</b>	0.0556	0.094	0.593	0.554	-0.130	0.241
<b>carwidth</b>	0.1912	0.098	1.942	0.054	-0.004	0.386
<b>carheight</b>	0.0688	0.051	1.356	0.178	-0.032	0.169
<b>curbweight</b>	0.2768	0.138	2.011	0.046	0.004	0.549
<b>enginetype</b>	0.0420	0.041	1.035	0.303	-0.038	0.122
<b>enginesize</b>	0.4369	0.146	2.997	0.003	0.148	0.725
<b>boreratio</b>	-0.0255	0.051	-0.501	0.617	-0.126	0.075
<b>stroke</b>	-0.1112	0.054	-2.073	0.040	-0.217	-0.005
<b>compressionratio</b>	0.0365	0.047	0.778	0.438	-0.056	0.130
<b>horsepower</b>	0.2358	0.137	1.724	0.087	-0.035	0.507
<b>peakrpm</b>	0.0570	0.053	1.074	0.285	-0.048	0.162
<b>citympg</b>	-0.0622	0.196	-0.317	0.751	-0.450	0.326
<b>highwaympg</b>	0.0939	0.192	0.490	0.625	-0.286	0.473
<b>Omnibus:</b>	37.735	<b>Durbin-Watson:</b>	1.819			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	155.443			
<b>Skew:</b>	0.851	<b>Prob(JB):</b>	1.76e-34			
<b>Kurtosis:</b>	7.816	<b>Cond. No.</b>	90.0			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [29]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[29]:

	Features	VIF
18	highwaympg	143.64
17	citympg	131.82
9	curbweight	78.98
7	carwidth	54.81
6	wheelbase	42.19
11	enginesize	39.92
15	horsepower	33.82
3	carbody	26.73
13	stroke	18.45
8	carheight	17.50
12	boreratio	15.97
4	drivewheel	14.30
16	peakrpm	11.28
10	enginetype	8.93
0	CarName	5.80
2	doornumber	3.87
14	compressionratio	3.85
1	aspiration	2.36
5	enginelocation	1.36

**Model 7 - Dropping the citympg predictor as it has the highest P value (insignificane) and VIF**

In [30]:

```
X = X.drop('citympg',1)
X_train_lm = sm.add_constant(X)
lr_7 = sm.OLS(y_train, X_train_lm).fit()
lr_7.summary()
```

Out[30]:

## OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.884
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.867
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	52.47
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	3.64e-49
<b>Time:</b>	10:34:07	<b>Log-Likelihood:</b>	170.94
<b>No. Observations:</b>	143	<b>AIC:</b>	-303.9
<b>Df Residuals:</b>	124	<b>BIC:</b>	-247.6
<b>Df Model:</b>	18		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.1208	0.089	-1.365	0.175	-0.296	0.054
<b>CarName</b>	-0.1242	0.027	-4.520	0.000	-0.179	-0.070
<b>aspiration</b>	0.0009	0.023	0.040	0.968	-0.045	0.046
<b>doornumber</b>	-0.0075	0.020	-0.377	0.707	-0.047	0.032
<b>carbody</b>	-0.0748	0.050	-1.508	0.134	-0.173	0.023
<b>drivewheel</b>	0.0619	0.036	1.730	0.086	-0.009	0.133
<b>enginelocation</b>	0.3226	0.092	3.521	0.001	0.141	0.504
<b>wheelbase</b>	0.0526	0.093	0.566	0.572	-0.131	0.236
<b>carwidth</b>	0.1921	0.098	1.958	0.052	-0.002	0.386
<b>carheight</b>	0.0677	0.050	1.342	0.182	-0.032	0.168
<b>curbweight</b>	0.2799	0.137	2.046	0.043	0.009	0.551
<b>enginetype</b>	0.0428	0.040	1.061	0.291	-0.037	0.123
<b>enginesize</b>	0.4219	0.137	3.070	0.003	0.150	0.694
<b>boreratio</b>	-0.0220	0.049	-0.445	0.657	-0.120	0.076
<b>stroke</b>	-0.1057	0.051	-2.090	0.039	-0.206	-0.006
<b>compressionratio</b>	0.0344	0.046	0.742	0.459	-0.057	0.126
<b>horsepower</b>	0.2502	0.129	1.945	0.054	-0.004	0.505
<b>peakrpm</b>	0.0554	0.053	1.052	0.295	-0.049	0.160
<b>highwaympg</b>	0.0396	0.087	0.458	0.648	-0.132	0.211

<b>Omnibus:</b>	38.156	<b>Durbin-Watson:</b>	1.819
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	157.707
<b>Skew:</b>	0.862	<b>Prob(JB):</b>	5.68e-35
<b>Kurtosis:</b>	7.847	<b>Cond. No.</b>	58.2

## Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



In [31]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[31]:

	Features	VIF
9	curbweight	78.96
7	carwidth	54.79
6	wheelbase	41.86
11	enginesize	35.81
15	horsepower	29.88
3	carbody	26.69
8	carheight	17.34
13	stroke	16.85
12	boreratio	15.44
17	highwaympg	15.24
4	drivewheel	14.25
16	peakrpm	11.00
10	enginetype	8.93
0	CarName	5.77
14	compressionratio	3.82
2	doornumber	3.74
1	aspiration	2.18
5	enginelocation	1.36

**Model 8 - Dropping the aspiration predictor as it has the highest P value (insignificane) and VIF**

In [32]:

```
X = X.drop('aspiration',1)
X_train_lm = sm.add_constant(X)
lr_8 = sm.OLS(y_train, X_train_lm).fit()
lr_8.summary()
```

Out[32]:

## OLS Regression Results

Dep. Variable:	price		R-squared:		0.884	
Model:	OLS		Adj. R-squared:		0.868	
Method:	Least Squares		F-statistic:		56.01	
Date:	Wed, 02 Oct 2019		Prob (F-statistic):		4.78e-50	
Time:	10:34:07		Log-Likelihood:		170.94	
No. Observations:	143		AIC:		-305.9	
Df Residuals:	125		BIC:		-252.6	
Df Model:	17					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-0.1203	0.087	-1.378	0.171	-0.293	0.053
CarName	-0.1243	0.027	-4.564	0.000	-0.178	-0.070
doornumber	-0.0075	0.020	-0.377	0.707	-0.047	0.032
carbody	-0.0748	0.049	-1.514	0.133	-0.173	0.023
drivewheel	0.0615	0.034	1.804	0.074	-0.006	0.129
enginelocation	0.3227	0.091	3.536	0.001	0.142	0.503
wheelbase	0.0530	0.092	0.578	0.565	-0.129	0.235
carwidth	0.1918	0.097	1.969	0.051	-0.001	0.385
carheight	0.0675	0.050	1.352	0.179	-0.031	0.166
curbweight	0.2806	0.135	2.078	0.040	0.013	0.548
enginetype	0.0427	0.040	1.065	0.289	-0.037	0.122
enginesize	0.4193	0.120	3.485	0.001	0.181	0.657
boreratio	-0.0220	0.049	-0.446	0.656	-0.119	0.075
stroke	-0.1053	0.049	-2.132	0.035	-0.203	-0.008
compressionratio	0.0350	0.044	0.802	0.424	-0.051	0.121
horsepower	0.2520	0.120	2.107	0.037	0.015	0.489
peakrpm	0.0549	0.051	1.080	0.282	-0.046	0.155
highwaympg	0.0392	0.086	0.458	0.648	-0.130	0.209
Omnibus:	38.092	Durbin-Watson:	1.818			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	157.385			
Skew:	0.860	Prob(JB):	6.67e-35			
Kurtosis:	7.843	Cond. No.	54.8			

## Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [33]:

```

vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

Out[33]:

	Features	VIF
8	curbweight	75.77
6	carwidth	54.53
5	wheelbase	41.35
10	enginesize	27.51
2	carbody	26.67
14	horsepower	26.22
7	carheight	17.19
12	stroke	15.80
11	boreratio	15.36
16	highwaympg	15.24
3	drivewheel	13.18
15	peakrpm	10.50
9	enginetype	8.93
0	CarName	5.74
1	doornumber	3.71
13	compressionratio	3.48
4	enginelocation	1.36

**Model 9 - Dropping the highwaympg predictor as it has the highest P value (insignificane) and VIF**

In [34]:

```
X = X.drop('highwaympg',1)
X_train_lm = sm.add_constant(X)
lr_9 = sm.OLS(y_train, X_train_lm).fit()
lr_9.summary()
```

Out[34]:

## OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.884
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.869
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	59.87
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	6.73e-51
<b>Time:</b>	10:34:07	<b>Log-Likelihood:</b>	170.82
<b>No. Observations:</b>	143	<b>AIC:</b>	-307.6
<b>Df Residuals:</b>	126	<b>BIC:</b>	-257.3
<b>Df Model:</b>	16		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0922	0.062	-1.489	0.139	-0.215	0.030
<b>CarName</b>	-0.1241	0.027	-4.572	0.000	-0.178	-0.070
<b>doornumber</b>	-0.0078	0.020	-0.395	0.694	-0.047	0.031
<b>carbody</b>	-0.0714	0.049	-1.467	0.145	-0.168	0.025
<b>drivewheel</b>	0.0625	0.034	1.842	0.068	-0.005	0.130
<b>enginelocation</b>	0.3216	0.091	3.537	0.001	0.142	0.502
<b>wheelbase</b>	0.0470	0.091	0.519	0.605	-0.132	0.226
<b>carwidth</b>	0.1959	0.097	2.026	0.045	0.005	0.387
<b>carheight</b>	0.0682	0.050	1.372	0.172	-0.030	0.167
<b>curbweight</b>	0.2444	0.109	2.240	0.027	0.028	0.460
<b>enginetype</b>	0.0371	0.038	0.975	0.332	-0.038	0.113
<b>engineize</b>	0.4359	0.114	3.812	0.000	0.210	0.662
<b>boreratio</b>	-0.0243	0.049	-0.497	0.620	-0.121	0.072
<b>stroke</b>	-0.1054	0.049	-2.142	0.034	-0.203	-0.008
<b>compressionratio</b>	0.0455	0.037	1.230	0.221	-0.028	0.119
<b>horsepower</b>	0.2440	0.118	2.069	0.041	0.011	0.477
<b>peakrpm</b>	0.0513	0.050	1.024	0.308	-0.048	0.150

<b>Omnibus:</b>	39.127	<b>Durbin-Watson:</b>	1.825
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	162.954
<b>Skew:</b>	0.886	<b>Prob(JB):</b>	4.12e-36
<b>Kurtosis:</b>	7.920	<b>Cond. No.</b>	49.3

## Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [35]:

```

vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif

```

Out[35]:

	Features	VIF
8	curbweight	58.78
6	carwidth	52.85
5	wheelbase	38.52
14	horsepower	24.44
10	enginesize	24.11
2	carbody	22.15
7	carheight	16.35
11	boreratio	15.18
12	stroke	14.06
3	drivewheel	12.37
15	peakrpm	10.32
9	enginetype	8.92
0	CarName	5.49
1	doornumber	3.49
13	compressionratio	2.96
4	enginelocation	1.35

**Model 10 - Dropping the boreratio predictor as it has the highest P value (insignificane) and VIF**

In [36]:

```
X = X.drop('boreratio',1)
X_train_lm = sm.add_constant(X)
lr_10 = sm.OLS(y_train, X_train_lm).fit()
lr_10.summary()
```



Out[36]:

## OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.884
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.870
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	64.23
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	9.29e-52
<b>Time:</b>	10:34:07	<b>Log-Likelihood:</b>	170.68
<b>No. Observations:</b>	143	<b>AIC:</b>	-309.4
<b>Df Residuals:</b>	127	<b>BIC:</b>	-262.0
<b>Df Model:</b>	15		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.1002	0.060	-1.678	0.096	-0.218	0.018
<b>CarName</b>	-0.1282	0.026	-4.980	0.000	-0.179	-0.077
<b>doornumber</b>	-0.0080	0.020	-0.403	0.688	-0.047	0.031
<b>carbody</b>	-0.0725	0.048	-1.496	0.137	-0.169	0.023
<b>drivewheel</b>	0.0601	0.033	1.794	0.075	-0.006	0.126
<b>enginelocation</b>	0.3150	0.090	3.512	0.001	0.137	0.492
<b>wheelbase</b>	0.0495	0.090	0.549	0.584	-0.129	0.228
<b>carwidth</b>	0.1862	0.094	1.972	0.051	-0.001	0.373
<b>carheight</b>	0.0699	0.049	1.413	0.160	-0.028	0.168
<b>curbweight</b>	0.2357	0.107	2.195	0.030	0.023	0.448
<b>enginetype</b>	0.0372	0.038	0.979	0.330	-0.038	0.112
<b>engine size</b>	0.4385	0.114	3.850	0.000	0.213	0.664
<b>stroke</b>	-0.0983	0.047	-2.094	0.038	-0.191	-0.005
<b>compressionratio</b>	0.0471	0.037	1.283	0.202	-0.026	0.120
<b>horsepower</b>	0.2424	0.118	2.062	0.041	0.010	0.475
<b>peakrpm</b>	0.0568	0.049	1.168	0.245	-0.039	0.153

<b>Omnibus:</b>	37.208	<b>Durbin-Watson:</b>	1.820
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	146.946
<b>Skew:</b>	0.853	<b>Prob(JB):</b>	1.23e-32
<b>Kurtosis:</b>	7.664	<b>Cond. No.</b>	47.7

## Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [37]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[37]:

	Features	VIF
8	curbweight	57.32
6	carwidth	49.41
5	wheelbase	37.95
13	horsepower	24.42
10	enginesize	24.11
2	carbody	21.44
7	carheight	16.35
11	stroke	13.49
3	drivewheel	11.69
14	peakrpm	10.09
9	enginetype	8.87
0	CarName	4.58
1	doornumber	3.43
12	compressionratio	2.93
4	enginelocation	1.33

**Model 11 - Dropping the carwidth predictor as it has the highest VIF value**

In [38]:

```
X = X.drop('carwidth',1)
X_train_lm = sm.add_constant(X)
lr_11 = sm.OLS(y_train, X_train_lm).fit()
lr_11.summary()
```

Out[38]:

## OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.880
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.867
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	67.02
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	7.36e-52
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	168.52
<b>No. Observations:</b>	143	<b>AIC:</b>	-307.0
<b>Df Residuals:</b>	128	<b>BIC:</b>	-262.6
<b>Df Model:</b>	14		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0742	0.059	-1.261	0.210	-0.191	0.042
<b>CarName</b>	-0.1305	0.026	-5.018	0.000	-0.182	-0.079
<b>doornumber</b>	-0.0128	0.020	-0.647	0.519	-0.052	0.026
<b>carbody</b>	-0.0980	0.047	-2.074	0.040	-0.192	-0.005
<b>drivewheel</b>	0.0441	0.033	1.342	0.182	-0.021	0.109
<b>enginelocation</b>	0.2856	0.089	3.194	0.002	0.109	0.463
<b>wheelbase</b>	0.1456	0.077	1.897	0.060	-0.006	0.298
<b>carheight</b>	0.0652	0.050	1.304	0.195	-0.034	0.164
<b>curbweight</b>	0.2720	0.107	2.543	0.012	0.060	0.484
<b>enginetype</b>	0.0587	0.037	1.597	0.113	-0.014	0.132
<b>enginesize</b>	0.4367	0.115	3.792	0.000	0.209	0.664
<b>stroke</b>	-0.0941	0.047	-1.984	0.049	-0.188	-0.000
<b>compressionratio</b>	0.0662	0.036	1.847	0.067	-0.005	0.137
<b>horsepower</b>	0.3232	0.111	2.901	0.004	0.103	0.544
<b>peakrpm</b>	0.0677	0.049	1.384	0.169	-0.029	0.164

<b>Omnibus:</b>	29.586	<b>Durbin-Watson:</b>	1.847
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	110.726
<b>Skew:</b>	0.652	<b>Prob(JB):</b>	9.04e-25
<b>Kurtosis:</b>	7.109	<b>Cond. No.</b>	43.3

## Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [39]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[39]:

	Features	VIF
7	curbweight	55.51
5	wheelbase	27.95
9	enginesize	24.07
12	horsepower	22.00
2	carbody	20.65
6	carheight	16.35
10	stroke	13.17
3	drivewheel	11.30
13	peakrpm	9.71
8	enginetype	7.57
0	CarName	4.57
1	doornumber	3.43
11	compressionratio	2.73
4	enginelocation	1.29

## Model 12 - Dropping the wheelbase predictor as it has the highest VIF value

In [40]:

```
X = X.drop('wheelbase',1)
X_train_lm = sm.add_constant(X)
lr_12 = sm.OLS(y_train, X_train_lm).fit()
lr_12.summary()
```

Out[40]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.877
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.864
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	70.48
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	4.94e-52
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	166.54
<b>No. Observations:</b>	143	<b>AIC:</b>	-305.1
<b>Df Residuals:</b>	129	<b>BIC:</b>	-263.6
<b>Df Model:</b>	13		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0924	0.059	-1.574	0.118	-0.208	0.024
<b>CarName</b>	-0.1291	0.026	-4.917	0.000	-0.181	-0.077
<b>doornumber</b>	-0.0177	0.020	-0.890	0.375	-0.057	0.022
<b>carbody</b>	-0.0847	0.047	-1.794	0.075	-0.178	0.009
<b>drivewheel</b>	0.0650	0.031	2.079	0.040	0.003	0.127
<b>enginelocation</b>	0.2560	0.089	2.878	0.005	0.080	0.432
<b>carheight</b>	0.0964	0.048	2.024	0.045	0.002	0.191
<b>curbweight</b>	0.3783	0.092	4.113	0.000	0.196	0.560
<b>enginetype</b>	0.0541	0.037	1.458	0.147	-0.019	0.127
<b>engine size</b>	0.4488	0.116	3.865	0.000	0.219	0.679
<b>stroke</b>	-0.0840	0.048	-1.766	0.080	-0.178	0.010
<b>compressionratio</b>	0.0655	0.036	1.810	0.073	-0.006	0.137
<b>horsepower</b>	0.2718	0.109	2.490	0.014	0.056	0.488
<b>peakrpm</b>	0.0770	0.049	1.568	0.119	-0.020	0.174

<b>Omnibus:</b>	25.502	<b>Durbin-Watson:</b>	1.848
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	97.498
<b>Skew:</b>	0.514	<b>Prob(JB):</b>	6.74e-22
<b>Kurtosis:</b>	6.912	<b>Cond. No.</b>	40.1

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [41]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[41]:

	Features	VIF
6	curbweight	40.27
8	enginesize	24.06
11	horsepower	20.91
2	carbody	20.48
5	carheight	14.86
9	stroke	13.14
3	drivewheel	10.25
12	peakrpm	9.69
7	enginetype	7.41
0	CarName	4.56
1	doornumber	3.28
10	compressionratio	2.73
4	enginelocation	1.25

## Model 13 - Dropping the carheight predictor as it has the highest VIF value

In [42]:

```
X = X.drop('carheight',1)
X_train_lm = sm.add_constant(X)
lr_13 = sm.OLS(y_train, X_train_lm).fit()
lr_13.summary()
```

Out[42]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.873
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.861
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	74.25
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	4.10e-52
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	164.31
<b>No. Observations:</b>	143	<b>AIC:</b>	-302.6
<b>Df Residuals:</b>	130	<b>BIC:</b>	-264.1
<b>Df Model:</b>	12		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0654	0.058	-1.132	0.260	-0.180	0.049
<b>CarName</b>	-0.1226	0.026	-4.647	0.000	-0.175	-0.070
<b>doornumber</b>	-0.0188	0.020	-0.939	0.350	-0.059	0.021
<b>carbody</b>	-0.0483	0.044	-1.093	0.276	-0.136	0.039
<b>drivewheel</b>	0.0585	0.031	1.860	0.065	-0.004	0.121
<b>enginelocation</b>	0.2907	0.088	3.292	0.001	0.116	0.465
<b>curbweight</b>	0.4590	0.084	5.472	0.000	0.293	0.625
<b>enginetype</b>	0.0413	0.037	1.116	0.266	-0.032	0.114
<b>enginesize</b>	0.4227	0.117	3.620	0.000	0.192	0.654
<b>stroke</b>	-0.0887	0.048	-1.844	0.068	-0.184	0.006
<b>compressionratio</b>	0.0641	0.037	1.752	0.082	-0.008	0.137
<b>horsepower</b>	0.2008	0.105	1.920	0.057	-0.006	0.408
<b>peakrpm</b>	0.0712	0.050	1.435	0.154	-0.027	0.169

<b>Omnibus:</b>	28.781	<b>Durbin-Watson:</b>	1.808
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	101.345
<b>Skew:</b>	0.652	<b>Prob(JB):</b>	9.85e-23
<b>Kurtosis:</b>	6.912	<b>Cond. No.</b>	38.5

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



In [43]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[43]:

	Features	VIF
5	curbweight	32.65
7	enginesize	23.94
10	horsepower	18.16
2	carbody	14.68
8	stroke	13.11
3	drivewheel	10.23
11	peakrpm	9.69
6	enginetype	7.35
0	CarName	4.31
1	doornumber	3.26
9	compressionratio	2.73
4	enginelocation	1.20

## Model 14 - Dropping the enginetype predictor as it has the high P Value

In [44]:

```
X = X.drop('enginetype',1)
X_train_lm = sm.add_constant(X)
lr_14 = sm.OLS(y_train, X_train_lm).fit()
lr_14.summary()
```

Out[44]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.871
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.861
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	80.73
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	8.22e-53
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	163.62
<b>No. Observations:</b>	143	<b>AIC:</b>	-303.2
<b>Df Residuals:</b>	131	<b>BIC:</b>	-267.7
<b>Df Model:</b>	11		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0402	0.053	-0.754	0.452	-0.146	0.065
<b>CarName</b>	-0.1248	0.026	-4.741	0.000	-0.177	-0.073
<b>doornumber</b>	-0.0165	0.020	-0.826	0.410	-0.056	0.023
<b>carbody</b>	-0.0443	0.044	-1.005	0.317	-0.131	0.043
<b>drivewheel</b>	0.0542	0.031	1.734	0.085	-0.008	0.116
<b>enginelocation</b>	0.2978	0.088	3.378	0.001	0.123	0.472
<b>curbweight</b>	0.4602	0.084	5.481	0.000	0.294	0.626
<b>enginesize</b>	0.4231	0.117	3.620	0.000	0.192	0.654
<b>stroke</b>	-0.0993	0.047	-2.103	0.037	-0.193	-0.006
<b>compressionratio</b>	0.0654	0.037	1.786	0.076	-0.007	0.138
<b>horsepower</b>	0.1995	0.105	1.906	0.059	-0.008	0.407
<b>peakrpm</b>	0.0731	0.050	1.472	0.143	-0.025	0.171

<b>Omnibus:</b>	28.581	<b>Durbin-Watson:</b>	1.820
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	90.054
<b>Skew:</b>	0.691	<b>Prob(JB):</b>	2.79e-20
<b>Kurtosis:</b>	6.634	<b>Cond. No.</b>	37.1

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [45]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[45]:

	Features	VIF
5	curbweight	32.60
6	enginesize	23.79
9	horsepower	17.95
7	stroke	13.09
2	carbody	12.50
3	drivewheel	10.22
10	peakrpm	9.37
0	CarName	4.25
1	doornumber	2.94
8	compressionratio	2.73
4	enginelocation	1.20

## Model 15 - Dropping the curbweight predictor as it has the high VIF Value

In [46]:

```
X = X.drop('curbweight',1)
X_train_lm = sm.add_constant(X)
lr_15 = sm.OLS(y_train, X_train_lm).fit()
lr_15.summary()
```

Out[46]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.842
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.830
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	70.33
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	6.10e-48
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	148.86
<b>No. Observations:</b>	143	<b>AIC:</b>	-275.7
<b>Df Residuals:</b>	132	<b>BIC:</b>	-243.1
<b>Df Model:</b>	10		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0191	0.059	-0.326	0.745	-0.135	0.097
<b>CarName</b>	-0.1180	0.029	-4.064	0.000	-0.175	-0.061
<b>doornumber</b>	-0.0473	0.021	-2.234	0.027	-0.089	-0.005
<b>carbody</b>	-0.0074	0.048	-0.154	0.878	-0.103	0.088
<b>drivewheel</b>	0.0986	0.033	2.959	0.004	0.033	0.165
<b>enginelocation</b>	0.1909	0.095	2.010	0.046	0.003	0.379
<b>engine size</b>	0.6842	0.118	5.804	0.000	0.451	0.917
<b>stroke</b>	-0.0836	0.052	-1.606	0.111	-0.187	0.019
<b>compressionratio</b>	0.1165	0.039	2.977	0.003	0.039	0.194
<b>horsepower</b>	0.4464	0.104	4.278	0.000	0.240	0.653
<b>peakrpm</b>	0.0402	0.054	0.739	0.461	-0.067	0.148

<b>Omnibus:</b>	18.335	<b>Durbin-Watson:</b>	1.848
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	37.844
<b>Skew:</b>	0.538	<b>Prob(JB):</b>	6.06e-09
<b>Kurtosis:</b>	5.279	<b>Cond. No.</b>	36.0

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [47]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[47]:

	Features	VIF
5	enginesize	19.38
8	horsepower	14.68
6	stroke	12.97
2	carbody	11.67
3	drivewheel	9.40
9	peakrpm	9.27
0	CarName	4.22
1	doornumber	2.70
7	compressionratio	2.55
4	enginelocation	1.14

## Model 16 - Dropping the carbody predictor as it has the high P Value

In [48]:

```
X = X.drop('carbody',1)
X_train_lm = sm.add_constant(X)
lr_16 = sm.OLS(y_train, X_train_lm).fit()
lr_16.summary()
```

Out[48]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.842
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.831
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	78.72
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	6.74e-49
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	148.85
<b>No. Observations:</b>	143	<b>AIC:</b>	-277.7
<b>Df Residuals:</b>	133	<b>BIC:</b>	-248.1
<b>Df Model:</b>	9		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0250	0.044	-0.566	0.572	-0.113	0.062
<b>CarName</b>	-0.1180	0.029	-4.080	0.000	-0.175	-0.061
<b>doornumber</b>	-0.0453	0.017	-2.732	0.007	-0.078	-0.012
<b>drivewheel</b>	0.0992	0.033	3.009	0.003	0.034	0.164
<b>enginelocation</b>	0.1932	0.093	2.067	0.041	0.008	0.378
<b>enginesize</b>	0.6841	0.117	5.825	0.000	0.452	0.916
<b>stroke</b>	-0.0832	0.052	-1.606	0.111	-0.186	0.019
<b>compressionratio</b>	0.1160	0.039	2.986	0.003	0.039	0.193
<b>horsepower</b>	0.4457	0.104	4.291	0.000	0.240	0.651
<b>peakrpm</b>	0.0396	0.054	0.733	0.465	-0.067	0.147

<b>Omnibus:</b>	18.131	<b>Durbin-Watson:</b>	1.847
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	37.161
<b>Skew:</b>	0.534	<b>Prob(JB):</b>	8.52e-09
<b>Kurtosis:</b>	5.257	<b>Cond. No.</b>	33.3

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [49]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[49]:

	Features	VIF
4	enginesize	18.41
7	horsepower	14.33
5	stroke	11.78
2	drivewheel	9.37
8	peakrpm	7.58
0	CarName	3.44
6	compressionratio	2.50
1	doornumber	2.15
3	enginelocation	1.11

## Model 17 - Dropping the peakrpm predictor as it has the high P Value and VIF Value

In [50]:

```
X = X.drop('peakrpm',1)
X_train_lm = sm.add_constant(X)
lr_17 = sm.OLS(y_train, X_train_lm).fit()
lr_17.summary()
```

Out[50]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.841
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.832
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	88.80
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	9.01e-50
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	148.56
<b>No. Observations:</b>	143	<b>AIC:</b>	-279.1
<b>Df Residuals:</b>	134	<b>BIC:</b>	-252.5
<b>Df Model:</b>	8		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0078	0.037	-0.209	0.835	-0.082	0.066
<b>CarName</b>	-0.1218	0.028	-4.288	0.000	-0.178	-0.066
<b>doornumber</b>	-0.0441	0.016	-2.679	0.008	-0.077	-0.012
<b>drivewheel</b>	0.1010	0.033	3.077	0.003	0.036	0.166
<b>enginelocation</b>	0.2028	0.092	2.196	0.030	0.020	0.385
<b>enginesize</b>	0.6361	0.097	6.539	0.000	0.444	0.828
<b>stroke</b>	-0.0759	0.051	-1.495	0.137	-0.176	0.024
<b>compressionratio</b>	0.1071	0.037	2.907	0.004	0.034	0.180
<b>horsepower</b>	0.4818	0.091	5.278	0.000	0.301	0.662

<b>Omnibus:</b>	19.584	<b>Durbin-Watson:</b>	1.843
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	39.832
<b>Skew:</b>	0.586	<b>Prob(JB):</b>	2.24e-09
<b>Kurtosis:</b>	5.305	<b>Cond. No.</b>	26.6

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



In [51]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[51]:

	Features	VIF
4	enginesize	14.16
7	horsepower	12.00
2	drivewheel	8.84
5	stroke	7.55
0	CarName	3.31
6	compressionratio	2.28
1	doornumber	2.06
3	enginelocation	1.09

## Model 18 - Dropping the enginesize predictor as it has the high P Value and VIF Value

In [52]:

```
X = X.drop('enginesize',1)
X_train_lm = sm.add_constant(X)
lr_18 = sm.OLS(y_train, X_train_lm).fit()
lr_18.summary()
```

Out[52]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.791
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.780
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	72.84
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	9.67e-43
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	128.76
<b>No. Observations:</b>	143	<b>AIC:</b>	-241.5
<b>Df Residuals:</b>	135	<b>BIC:</b>	-217.8
<b>Df Model:</b>	7		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	0.0024	0.043	0.056	0.955	-0.082	0.087
<b>CarName</b>	-0.1464	0.032	-4.543	0.000	-0.210	-0.083
<b>doornumber</b>	-0.0741	0.018	-4.097	0.000	-0.110	-0.038
<b>drivewheel</b>	0.1334	0.037	3.590	0.000	0.060	0.207
<b>enginelocation</b>	0.1833	0.106	1.735	0.085	-0.026	0.392
<b>stroke</b>	-0.0218	0.057	-0.380	0.705	-0.135	0.092
<b>compressionratio</b>	0.1700	0.041	4.176	0.000	0.089	0.250
<b>horsepower</b>	0.9549	0.064	14.984	0.000	0.829	1.081

<b>Omnibus:</b>	34.934	<b>Durbin-Watson:</b>	1.640
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	72.555
<b>Skew:</b>	1.054	<b>Prob(JB):</b>	1.76e-16
<b>Kurtosis:</b>	5.780	<b>Cond. No.</b>	19.3

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [53]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[53]:

	Features	VIF
2	drivewheel	8.58
4	stroke	7.01
6	horsepower	4.47
0	CarName	3.25
5	compressionratio	2.13
1	doornumber	1.90
3	enginelocation	1.09

## Model 19 - Dropping the stroke predictor as it has the high P Value and VIF Value

In [54]:

```
X = X.drop('stroke',1)
X_train_lm = sm.add_constant(X)
lr_19 = sm.OLS(y_train, X_train_lm).fit()
lr_19.summary()
```

Out[54]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.790
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.781
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	85.50
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	1.07e-43
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	128.69
<b>No. Observations:</b>	143	<b>AIC:</b>	-243.4
<b>Df Residuals:</b>	136	<b>BIC:</b>	-222.6
<b>Df Model:</b>	6		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.0092	0.030	-0.307	0.759	-0.068	0.050
<b>CarName</b>	-0.1440	0.032	-4.570	0.000	-0.206	-0.082
<b>doornumber</b>	-0.0743	0.018	-4.121	0.000	-0.110	-0.039
<b>drivewheel</b>	0.1336	0.037	3.609	0.000	0.060	0.207
<b>enginelocation</b>	0.1881	0.105	1.800	0.074	-0.019	0.395
<b>compressionratio</b>	0.1664	0.039	4.214	0.000	0.088	0.245
<b>horsepower</b>	0.9516	0.063	15.119	0.000	0.827	1.076

<b>Omnibus:</b>	34.637	<b>Durbin-Watson:</b>	1.634
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	71.741
<b>Skew:</b>	1.046	<b>Prob(JB):</b>	2.64e-16
<b>Kurtosis:</b>	5.768	<b>Cond. No.</b>	17.9

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [55]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[55]:

	Features	VIF
2	drivewheel	7.68
5	horsepower	4.30
0	CarName	2.74
4	compressionratio	1.96
1	doornumber	1.76
3	enginelocation	1.05

## Model 20 - Dropping the drivewheel predictor as it has the high P Value and VIF Value

In [56]:

```
X = X.drop('drivewheel',1)
X_train_lm = sm.add_constant(X)
lr_20 = sm.OLS(y_train, X_train_lm).fit()
lr_20.summary()
```

Out[56]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.770
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.762
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	91.93
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	5.06e-42
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	122.15
<b>No. Observations:</b>	143	<b>AIC:</b>	-232.3
<b>Df Residuals:</b>	137	<b>BIC:</b>	-214.5
<b>Df Model:</b>	5		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	0.0390	0.028	1.398	0.164	-0.016	0.094
<b>CarName</b>	-0.1487	0.033	-4.529	0.000	-0.214	-0.084
<b>doornumber</b>	-0.0669	0.019	-3.582	0.000	-0.104	-0.030
<b>enginelocation</b>	0.1807	0.109	1.658	0.100	-0.035	0.396
<b>compressionratio</b>	0.2180	0.038	5.675	0.000	0.142	0.294
<b>horsepower</b>	1.0733	0.055	19.365	0.000	0.964	1.183

<b>Omnibus:</b>	36.381	<b>Durbin-Watson:</b>	1.559
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	78.682
<b>Skew:</b>	1.079	<b>Prob(JB):</b>	8.21e-18
<b>Kurtosis:</b>	5.923	<b>Cond. No.</b>	15.8

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [57]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[57]:

	Features	VIF
0	CarName	2.36
4	horsepower	2.32
1	doornumber	1.57
3	compressionratio	1.48
2	enginelocation	1.05

## Model 21 - Dropping the enginelocation predictor as it has the high P Value

In [58]:

```
X = X.drop('enginelocation',1)
X_train_lm = sm.add_constant(X)
lr_21 = sm.OLS(y_train, X_train_lm).fit()
lr_21.summary()
```

Out[58]:

OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.766
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.759
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	112.8
<b>Date:</b>	Wed, 02 Oct 2019	<b>Prob (F-statistic):</b>	1.72e-42
<b>Time:</b>	10:34:08	<b>Log-Likelihood:</b>	120.73
<b>No. Observations:</b>	143	<b>AIC:</b>	-231.5
<b>Df Residuals:</b>	138	<b>BIC:</b>	-216.6
<b>Df Model:</b>	4		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	0.0316	0.028	1.141	0.256	-0.023	0.086
<b>CarName</b>	-0.1443	0.033	-4.381	0.000	-0.209	-0.079
<b>doornumber</b>	-0.0640	0.019	-3.419	0.001	-0.101	-0.027
<b>compressionratio</b>	0.2203	0.039	5.703	0.000	0.144	0.297
<b>horsepower</b>	1.0936	0.054	20.104	0.000	0.986	1.201

<b>Omnibus:</b>	31.024	<b>Durbin-Watson:</b>	1.583
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	61.723
<b>Skew:</b>	0.954	<b>Prob(JB):</b>	3.95e-14
<b>Kurtosis:</b>	5.592	<b>Cond. No.</b>	8.19

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



In [59]:

```
vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[59]:


	Features	VIF
0	CarName	2.35
3	horsepower	2.25
1	doornumber	1.57
2	compressionratio	1.48

## Step 7 : Residual Analysis of the error data (train)

In [60]:

```
y_train_residual = lr_21.predict(X_train_lm)

# Plot the histogram of the error terms
fig = plt.figure()
sns.distplot((y_train - y_train_residual), bins = 20)
fig.suptitle('Error Terms', fontsize = 20)
plt.xlabel('Errors', fontsize = 18)
plt.show()
```



## Step 8 : Making the predictions using test data

In [61]:

```

master_df_test[num_columns_list] = scaler.fit_transform(master_df_test[num_columns_list])
master_df_test[non_num_columns_list] = scaler.fit_transform(master_df_test[non_num_columns_list])
master_df_test.describe()

```

Out[61]:

	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drivewh
count	62.000000	62.000000	62.000000	62.000000	62.000000	62.000000	62.000000
mean	0.583871	0.566532	0.887097	0.177419	0.435484	0.625000	0.701613
std	0.271724	0.311481	0.319058	0.385142	0.499868	0.225205	0.263343
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.400000	0.375000	1.000000	0.000000	0.000000	0.500000	0.500000
50%	0.600000	0.562500	1.000000	0.000000	0.000000	0.750000	0.500000
75%	0.800000	0.833333	1.000000	0.000000	1.000000	0.750000	1.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 25 columns

In [62]:

```

## Dividing between x_test and y_test

test_columns = list(master_df_test.columns)

y_test = master_df_test['price']
test_columns.remove('price')
x_test = master_df_test[test_columns]

```

In [63]:

```
y_test.head()
```

Out[63]:

```

car_ID
161    0.058474
187    0.077398
60     0.086148
166    0.097473
141    0.055099
Name: price, dtype: float64

```

In [64]:

```
x_test.head()
```

Out[64]:

	Insuranceriskfactor	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
car_ID							
161	0.4	0.833333	1.0	0.0	0.0	0.75	0.5
187	0.8	0.958333	1.0	0.0	0.0	0.75	0.5
60	0.6	0.375000	1.0	0.0	1.0	0.50	0.5
166	0.6	0.833333	1.0	0.0	1.0	0.75	1.0
141	0.8	0.791667	1.0	0.0	1.0	0.50	0.0

5 rows × 24 columns

In [65]:

```
## add constant variable to x_test
x_test_sm = sm.add_constant(x_test)

## Keeping only the relevant predictors from the final lr_20 model
x_test_pred = x_test_sm[['const', 'CarName', 'doornumber', 'compressionratio', 'horsepower']]

## Making the predictions on the test set
y_test_pred = lr_21.predict(x_test_pred)
```

In [66]:

```
y_test_pred.head()
```

Out[66]:


```
car_ID
161    0.065901
187    0.153693
60     0.161308
166    0.303734
141    0.024957
dtype: float64
```

## Step 9 : Model Evaluation

In [67]:

```
# Plotting y_test and y_pred to understand the spread

fig = plt.figure()
plt.scatter(y_test, y_test_pred)
fig.suptitle('y_test vs y_test_pred', fontsize = 20)
plt.xlabel('y_test', fontsize = 18)
plt.ylabel('y_test_pred', fontsize = 16)
plt.show()
```



## Step 10 : Calculate the R-square value

In [68]:

```
from sklearn.metrics import r2_score
r2_score(y_test, y_test_pred)
```

Out[68]:

0.07040807219687462

In [ ]: